

## THE UGSORT ALGORITHM

Tree, Ian J.

unaffiliated researcher

### Author Note

Ian J. Tree, unaffiliated researcher

Eindhoven, the Netherlands

Email: [ian.tree@acm.org](mailto:ian.tree@acm.org)

### Abstract

This paper describes the novel UGSort merge sorting algorithm. The algorithm is simple, elegant, flexible and easy to implement.

*Keywords:* sort, algorithm, merge

## The UGSort Algorithm

The UGSort algorithm (Unexpectedly Good Sort) is a simple, elegant and flexible algorithm. It is easy to implement in different settings and is efficient.

Revised: 11/09/2023 – Added binary search of partitions for new key.

Revised: 05/11/2023 – Pre-emptive merge changed to tail suppression.

### The Algorithm

The algorithm is a merge sorting algorithm based on partitioning the input keys into an array of double ended queues (dequeues) and then merging the double ended queues into a single queue from which the results of the sort can be output.

Sorting a set of keys  $KS = \{k_1, k_2, \dots, k_n\}$  with  $n$  keys begins with initialising the first element of an array ( $S$ ) of double ended queues  $S_0$  with the first key available from the input set  $k_1$ , both the head and the tail of  $S_0$  will be set to  $k_1$ .

Each subsequent key in the input set  $k_i = k_2, \dots, k_n$  will be added to the partitions by testing against the head and tail keys in each of the available double ended queues in  $S$ . The following describes a linear search of the partitions for the first qualifying partition, see the later section “Binary Search” for a description of replacing the linear search with a binary search. If the key  $k_i$  is less than or equal to the head value of the queue, then the key is added to the head of the queue. If the key was not consumed by the head of the queue, then if it is greater than or equal to the tail value of the queue then, it is added to the tail of the queue. If the key was not consumed by any of the queues in  $S$  then a new queue is added to the end of the array initialised with the unused key  $k_i$ .

If a new queue was added to the array  $S$  then the number of elements in the array is checked against a computed maximum  $S_{\max}$  if it exceeds the maximum then a pre-emptive merge is performed. The value of  $S_{\max}$  is the greater of 100 or  $(\frac{\text{sqrt}(\frac{i}{2})}{2})$ , where  $i$  is the number of keys processed so far. The pre-emptive merge will reduce the number of queues in the

array  $S$  by suppressing the tail of the array by merging the last (highest index)  $1/10^{\text{th}}$  elements into a single queue.

After the last key  $k_n$  has been added to the queues in the array  $S$  a final merge is performed where all queues are merged into the first queue in the array  $S_0$ .

The input set  $KS$  is now sorted into ascending sequence in the queue  $S_0$ .

### Pseudocode

```

Sort Input:       $KS = \{k_1, k_2 \dots k_n\}$ 

Partition Array:   $S = \{\}$ 

Max Partitions:   $m = 100$ 

 $S_0 = \text{double\_ended\_queue}(k_1)$ 

ForEach  $k_i$  in  $k_2 \dots k_n$ 
    ForEach  $S_j$  in  $S$ 
        If  $k_i \leq S_j.\text{head}$ 
             $S_j.\text{head} = k_i$ 
             $\text{Key\_consumed} = \text{true}$ 
        Exit ForEach
    End If
    If  $k_i \geq S_j.\text{tail}$ 
         $S_j.\text{tail} = k_i$ 
         $\text{Key\_consumed} = \text{true}$ 
    Exit ForEach
End If
End ForEach

If ( $\neg \text{Key\_consumed}$ )
     $S_{j+1} = \text{double\_ended\_queue}(k_i)$ 

```

```

    If (j+1 > m)

        // Perform pre-emptive merge

        Merged = 0, Target = j/10

        While Merged < Target

            Merge  $S_{last}$  into  $S_{penultimate}$ 

            Delete  $S_{last}$ 

            Merged++

        End While

        NewMax = (sqrt(i/2)/2)

        If (NewMax > m) m = NewMax

    End If

End If

End ForEach

// Perform final merge

ForEach  $S_j$  in  $S_1...S_m$ 

    Merge  $S_j$  into  $S_0$ 

End ForEach

```

## Binary Search

The linear search of the partitions to locate the correct one to add a new key to may be replaced by a more efficient binary search. The linear search finds the first partition where the new key can be added to the head or the tail of the partition, proceeding from the lowest index partition  $S_0$  to the last partition  $S_m$ , the binary search must mirror this constraint.

For each new key  $k_i$  begin by testing against the boundary conditions for the complete array of partitions. If the new key is less than or equal to the low key of the first partition  $S_0$  then add the key to  $S_0$  as the new low key. If the new key is greater than or equal to the high

key of the first partition  $S_0$  then add the key to  $S_0$  as the new high key. If the key was not consumed by the first  $S_0$  partition then test to see if the key will not be consumed by the last partition  $S_m$  and therefore cause a new partition to be initialised with the key. If the key was not consumed by either of the boundary conditions then start at the mid-point of the array  $S_{m/2}$  with a delta of  $m/4$ . At each partition test the new key to see if it would qualify for addition to the current partition if so then test if it would also qualify for the preceding partition if so subtract the delta from the current partition array index and halve the delta. If the key did not qualify for addition to the preceding partition then add the key to the current partition as the new high or low key as appropriate. If the key did not qualify for the current partition then also test it against the following partition. If the key qualifies for the following partition then add it as the new high or low key to the following partition as appropriate. If it did not qualify for the following partition then add the delta to the partition array index and halve the delta. The search continues until the new key is consumed by a partition.

If a new queue was added to the array  $S$  then the number of elements in the array is checked against a computed maximum  $S_{\max}$  if it exceeds the maximum then a pre-emptive merge is performed. The value of  $S_{\max}$  is the greater of 100 or  $(\frac{\text{sqrt}(\frac{i}{2})}{2})$ , where  $i$  is the number of keys processed so far. The pre-emptive merge will reduce the number of queues in the array  $S$  by suppressing the tail of the array by merging the last (highest index)  $1/10^{\text{th}}$  elements into a single queue.

### Binary Search Pseudocode

This snippet of pseudocode replaces the linear search in the reference pseudocode, it is executed for each new key  $k_i$  in  $k_2...k_n$ .

If  $k_i \leq S_0.\text{head}$

$S_0.\text{head} = k_i$

```

Exit
End If
If  $k_i \geq S_0.\text{tail}$ 
     $S_0.\text{tail} = k_i$ 
    Exit
End If
If  $k_i > S_j.\text{head}$  and  $k_i < S_j.\text{tail}$ 
     $S_{j+1} = \text{double\_ended\_queue}(k_i)$ 
    If  $(j+1 > m)$ 
        // Perform pre-emptive merge
        Merged = 0, Target =  $j/10$ 
        While Merged < Target
            Merge  $S_{\text{last}}$  into  $S_{\text{penultimate}}$ 
            Delete  $S_{\text{last}}$ 
            Merged++
        End While
        NewMax =  $(\text{sqrt}(i/2)/2)$ 
        If (NewMax > m) m = NewMax
    End If
    Exit
End If
// Perform Binary Search
 $c = m/2$  // Midpoint of partition array
 $\text{delta} = m/4$  // Binary chop delta
Loop

```

```

    If  $k_i > S_c.\text{head}$  and  $k_i < S_c.\text{tail}$ 
        If  $k_i > S_{c-1}.\text{head}$  and  $k_i < S_{c-1}.\text{tail}$ 
             $c = c - \text{delta}$ 
        Else
             $S_{c-1}.\text{head} = k_i$  or  $S_{c-1}.\text{tail} = k_i$ 
            Exit
    Else
        If  $k_i > S_{c+1}.\text{head}$  and  $k_i < S_{c+1}.\text{tail}$ 
             $c = c + \text{delta}$ 
        Else
             $S_c.\text{head} = k_i$  or  $S_c.\text{tail} = k_i$ 
            Exit
    EndIf
EndLoop

```

### Key Stability

Key stability is the sorting property of preserving the input sequence for keys of equal value. The UGSort algorithm as presented is not stable. However, it can be made stable by making minor changes to the algorithm. The tests for addition to the head or tail of a queue must be changed from less than or equal and greater than or equal to less than and greater than. The merge processing must always give precedence to the leftmost (lowest index) queue when keys are equal. These changes will result in a slightly less efficient algorithm.

### Best and Worst Cases

The “Best-Case” input for the algorithm is an input set SK that is pre-sorted into descending sequence, this results in a single queue being used with all keys being added to



the head of the queue. Pre-sorting into ascending sequence offers the next best case input, it also uses a single queue but will cost an additional key comparison for each key.

The notional "Worst-Case" input for the algorithm is an input set SK that is pre-sorted into ascending sequence and then reordered by taking alternating keys from the head and the tail of the sorted queue. For an ordered set of keys  $KS = \{k_1, k_2, \dots, k_n\}$  the worst-case input would be  $\{k_1, k_n, k_2, k_{n-1}, \dots\}$ . Reversing the sequence of a "Worst-Case" set of keys results in a set that will also be a "Best-Case" input.

### **Sort Properties**

As noted in the previous section the UGSort algorithm is not stable but, may be transformed into a stable sort by small modifications. The algorithm is not capable of in-place sorting and is therefore an out-of-place algorithm.

It is possible to introduce parallel processing for the final merge phases as individual merge operations can be done in parallel, there is no additional co-ordination needed for parallel final merges.